

# Stakeholder Engagement Plan

Project: Iroko — AI-Powered Mental Health Support Chatbot | Team Iroko

Iroko is a fine-tuned conversational AI built to provide culturally appropriate, safety-conscious mental health support for a Nigerian context, fine-tuned from Google's Gemma 4 E2B model. Because the system engages people in emotionally sensitive, sometimes high-risk conversations, its design and deployment decisions carry real consequences for the people who use it and for the professionals who may be asked to trust, refer to, or work alongside it. This plan identifies two key stakeholder groups, the benefits and risks the project poses to each, an appropriate engagement method drawn from the Action Catalogue for each group, and how their feedback would shape the project going forward.

## 1. Stakeholder Groups

### 1.1 Stakeholder Group A — End Users / Help-Seekers

| Aspect                    | Description                                                                                                                                                                                                                                                                                                                                                      |
|---------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Who they are              | Individuals in Nigeria — largely young adults and students — who turn to a chatbot for emotional support, a listening ear, or guidance when facing stress, anxiety, or more serious distress, including moments of crisis. Many would be first-time users of any mental health support tool, often because of cost, stigma, or lack of access to in-person care. |
| Why they matter           | They are the direct, intended beneficiaries of the system and the ones exposed to its failures. Every design decision — tone, safety behavior, crisis handling — is ultimately judged by whether it helps or harms this group in a real conversation.                                                                                                            |
| What matters most to them | Feeling heard and not judged; responses that reflect their cultural and linguistic context rather than a generic Western tone; privacy and confidentiality; and — critically — being pointed to real, local help (e.g. MANI, 112) when a conversation moves beyond what a chatbot should handle alone.                                                           |

### 1.2 Stakeholder Group B — Mental Health Professionals & Crisis Support Organizations

| Aspect          | Description                                                                                                                                                                                                                                                                                                                                 |
|-----------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Who they are    | Licensed counselors, psychologists, and crisis-response organizations in Nigeria — such as the Mentally Aware Nigeria Initiative (MANI) and emergency responders reachable via 112 — who either provide the care a chatbot cannot, or would be the ones a distressed user is escalated to.                                                  |
| Why they matter | They hold the clinical expertise the team lacks, and their judgment determines whether the chatbot's safety behavior (refusals, crisis detection, resource referrals) is actually sound rather than just plausible-sounding. Their willingness to be named as a referral point also affects the credibility and safety of the whole system. |

| Aspect                    | Description                                                                                                                                                                                                                                                                                          |
|---------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| What matters most to them | That the chatbot does not misrepresent itself as a substitute for professional care; that it reliably and quickly routes at-risk users to real help rather than attempting to manage crises itself; and that referencing their organization does not create liability or reputational risk for them. |

## 2. Potential Benefits and Risks

### 2.1 End Users / Help-Seekers

| Category                                              | Details                                                                                                                                                                                                                                                                                                                      |
|-------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Benefits                                              | Free, private, always-available first point of contact for emotional support; lowers the barrier posed by cost and stigma; can normalize seeking help at all.                                                                                                                                                                |
| Harms / concerns                                      | A model that gives poor or inconsistent advice in a genuine crisis; false reassurance that discourages a user from seeking real professional help; a user forming a one-sided emotional reliance on the chatbot; data privacy concerns given the sensitivity of what is shared.                                              |
| Ethical / social / cultural / economic considerations | Duty of care to a vulnerable population; the cultural appropriateness of responses (e.g. avoiding US-centric framing or resources); equitable access for users without reliable internet or English fluency; the economic reality that free tooling can widen or narrow the mental-health access gap depending on execution. |

### 2.2 Mental Health Professionals & Crisis Support Organizations

| Category                                              | Details                                                                                                                                                                                                                                                                                      |
|-------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Benefits                                              | A triage layer that can direct more people toward real services and existing hotlines; visibility for organizations like MANI among a population that might not otherwise know they exist.                                                                                                   |
| Harms / concerns                                      | Being referenced by a chatbot without consent or review, and being blamed if a user has a bad outcome after being pointed to them; the chatbot inadvertently displacing or delaying professional care rather than supporting it.                                                             |
| Ethical / social / cultural / economic considerations | Professional and legal accountability boundaries between an AI tool and licensed care; informed consent from partner organizations before being named as a resource; the risk of an unregulated AI tool shaping public perception of what mental health support should look like in Nigeria. |

## 3. Engagement Method (Action Catalogue)

### 3.1 End Users — Participatory Design

**Why this method is appropriate:** Participatory Design brings the people who will actually use the system into the design process itself, rather than treating them as passive test subjects — appropriate given how much tone, phrasing, and cultural fit matter to this group.

**Why it suits this stakeholder group:** End users are numerous, geographically dispersed, and diverse in language and background; a lightweight, iterative method that can run in small groups (rather than a single formal committee) fits their availability and comfort level, especially given the sensitivity of the topic.

**How it would be conducted:** Small, voluntary, anonymized sessions (in person or via chat) where participants review sample conversations and flag responses that feel tone-deaf, unclear, or unhelpful; iterate on phrasing and crisis-resource wording together; and pilot-test the interface before wider release, with participation kept low-friction and fully optional given the personal nature of the subject matter.

### 3.2 Mental Health Professionals & Crisis Organizations — Deliberative Workshop

**Why this method is appropriate:** A Deliberative Workshop brings a smaller number of expert stakeholders together for structured, in-depth discussion and joint reasoning about trade-offs — suited to technical and clinical review rather than broad public input.

**Why it suits this stakeholder group:** Professionals bring specialized judgment that benefits from focused deliberation on a shared set of materials (e.g. real anonymized transcripts and the red-team prompt set), rather than a large, loosely structured forum.

**How it would be conducted:** A half-day workshop with clinicians and a MANI representative, presenting flagged failure cases and the current safety harness; working through scenarios together to agree on what correct crisis behavior looks like; and formally asking for their consent and preferred wording before their organization is named as a referral resource in the chatbot.

## 4. Influence on Project Design

| Area         | How stakeholder feedback would shape it                                                                                                                                                                                                                                   |
|--------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Dataset      | End-user sessions surface phrasing and scenarios that feel foreign or unnatural, informing what gets added or rewritten in the training data; professional review flags conversation examples that model clinically inappropriate advice, which are removed or corrected. |
| Model design | Clinician input on where the model is too permissive or too blunt directly informs safety harness thresholds (e.g. what counts as an escalation trigger) and refusal-training priorities.                                                                                 |
| Outputs      | User feedback on tone shapes response style guidelines; professional feedback shapes the exact wording of crisis-resource messages so they match how local services actually want to be described.                                                                        |
| Evaluation   | Failure cases surfaced by both groups feed directly into the red-team prompt set, expanding which categories of adversarial and edge-case prompts are tested before each release.                                                                                         |
| Deployment   | Professional sign-off (or lack of it) determines whether specific organizations can be named as referral resources, and user pilot feedback determines whether the tool is ready to move beyond internal testing to a wider release.                                      |

## Appendix — Team & Contributors

Team members (Team Iroko):

- Ololade Ogunleye
- Lawal Habib
- Stephanie Omolu
- Alemoh Rapheal